

Rahul Kothagundla

+91-8500157826 | rahulkothagundla2002@gmail.com | LinkedIn | GitHub | Medium

PROFESSIONAL SUMMARY

Data Engineer with 2 years designing and maintaining production data pipelines on cloud platforms. Delivered 70% performance improvements using Azure Databricks and Spark, maintained 99.9% data reliability across 50K+ daily records and built real-time streaming systems with Kafka handling 1,000+ events/sec. Experienced with ETL/ELT workflows, cloud data warehousing and pipeline orchestration. Looking to design scalable data infrastructure that powers analytics and business intelligence.

TECHNICAL SKILLS

Languages: Python, SQL

Data Processing: Apache Spark, PySpark, Apache Kafka, Spark Structured Streaming, ETL/ELT Pipelines

Cloud Platforms: Azure (Data Factory, Databricks, Synapse), AWS (S3, Redshift, Glue)

Databases & Storage: PostgreSQL, Redis, SQL Server, Data Lakes, Data Warehousing

DevOps & Tools: Docker, Docker Compose, CI/CD, Git, Pipeline Monitoring, Data Quality Validation

PROFESSIONAL EXPERIENCE

Tata Consultancy Services

Jun 2024 – Present

Systems Engineer

Hyderabad, India

- Re-architected a financial data processing pipeline using Python and Azure Databricks, migrating heavy transformations to a distributed Spark backend and reducing pipeline execution time by 70%.
- Designed an automated ETL monitoring framework with error detection, failure recovery and secure backups, maintaining 99.9% data integrity across 50K+ daily records ingested from multiple sources.
- Built end-to-end ETL/ELT pipelines with multi-stage data validation and transformation logic, converting raw Salesforce CRM data into structured analytical assets stored in cloud data warehouses.
- Monitored and tuned query performance on cloud data infrastructure, optimizing resource consumption and reducing operational costs. Earned 5-star performance rating.
- Authored technical documentation covering data architecture and pipeline debugging, reducing engineer onboarding from 7 weeks to 4 days (90% improvement).

Atomstate Technologies

Jun 2023 – Sep 2023

Machine Learning Engineer

Hyderabad, India

- Built an automated data ingestion and processing pipeline for an NLP platform, handling large volumes of unstructured text through extraction, transformation and loading into analytical storage.
- Developed data-intensive workflows on large-scale datasets for a real-time analytics platform, generating structured insights for strategic decision-making.

Zummit Infolabs

Nov 2022 – Mar 2023

Junior Data Scientist

Remote

- Built data transformation pipelines converting unstructured text from APIs and databases into structured analytical datasets, enabling queryable business intelligence for clients.

TECHNICAL PROJECTS

Real-Time Stock Market Streaming Platform | Python, Kafka, Spark, PostgreSQL, Redis, Docker

- Designed and deployed a real-time streaming data pipeline ingesting live market data with sub-30s latency, handling 1,000+ events/sec using Apache Kafka with fault-tolerant configuration.
- Implemented Spark Structured Streaming for real-time data transformation with exactly-once semantics, feeding processed data into PostgreSQL with composite indexes optimized for analytical queries.
- Built Redis caching layer reducing database load by 60%, containerized all services with Docker Compose including health checks, achieving 99.9% pipeline uptime with auto-recovery.

EDUCATION

VNR Vignana Jyothi Institute of Engineering & Technology

2020 – 2024

B.Tech in Computer Science & Engineering – Data Science

CGPA: 8.55/10

CERTIFICATIONS & ACHIEVEMENTS

AWS Cloud Quest: Generative AI Practitioner

Amazon ML Summer School 2023